



Databases SQL MySQL PostgreSQL PL/SQL MongoDB SQL Cheat Sheet SQL Interview Questions My!

Top 50 MongoDB Interview Questions with Answers for 2024

Last Updated : 28 Aug, 2024

Getting ready for a MongoDB interview? MongoDB is a popular NoSQL database known for its flexibility and scalability, making it a favorite among many developers and companies. Whether you're a beginner looking to land your first job or an experienced professional aiming to advance your career, being well-prepared for your interview is crucial.

This blog post compiles a comprehensive list of MongoDB interview questions that will help you understand what to expect and how to tackle them confidently. Dive in and start preparing for a successful MongoDB interview.

Table of Content

- [MongoDB Basic Interview Questions](#)
- [MongoDB Intermediate Interview Questions](#)
- [MongoDB Query Based Interview Questions](#)



MongoDB Basic Interview Questions

MongoDB Basic Interview Questions focus on the essential knowledge you need to get started with this NoSQL database. These questions usually cover the differences between SQL and NoSQL databases, how MongoDB stores data in documents, and basic operations like creating, reading, updating, and deleting data.

1. What is MongoDB, and How Does It Differ from Traditional SQL Databases?

MongoDB is a [NoSQL](#) database, which means it does not use the traditional table-based relational database structure. Instead, it uses a flexible, document-oriented data model that stores data in BSON (Binary JSON) format. Unlike SQL databases that use rows and columns, MongoDB stores data as [JSON](#)-like documents, making it easier to handle unstructured data and providing greater flexibility in terms of schema design.

2. Explain BSON and Its Significance in MongoDB.

[BSON](#) (Binary JSON) is a binary-encoded serialization format used by MongoDB to store documents. BSON extends JSON by adding support for data types such as dates and binary data, and it is designed to be efficient in both storage space and scan speed. The binary format allows MongoDB to be more efficient with data retrieval and storage compared to text-based JSON.



Fitten Code: Code Faster, Code Better

3. Describe the Structure of a MongoDB Document.

A MongoDB document is a set of key-value pairs, similar to a JSON object. Each key is a string, and the value can be a variety of data types including strings, numbers, arrays, nested documents, and more. For example:

```
{
  "_id": ObjectId("507f1f77bcf86cd799439011"),
  "name": "Alice",
  "age": 25,
  "address": {
    "street": "123 Main St",
    "city": "Anytown",
    "state": "CA"
  },
  "hobbies": ["reading", "cycling"]
}
```

4. What are Collections And Databases In MongoDB?

In MongoDB, a database is a container for collections, and a collection is a group of MongoDB documents. Collections are equivalent to tables in [SQL](#) databases, but they do not enforce a schema, allowing for flexibility in the structure of the documents they contain. For example, you might have a users collection in a mydatabase database.

5. How Does MongoDB Ensure High Availability and Scalability?

MongoDB ensures high availability and scalability through its features like replica sets and sharding. Replica sets provide redundancy and failover capabilities, ensuring that data is always available. Sharding distributes data across multiple servers, enabling horizontal scalability to handle large volumes of data and high traffic loads.

6. Explain the Concept of Replica Sets in MongoDB.

A replica set in MongoDB is a group of mongod instances that maintain the same data set. A replica set consists of a primary node and multiple secondary nodes. The primary node receives all write operations, while secondary nodes replicate the primary's data and can serve read operations. If the primary node fails, an automatic election process selects a new primary to maintain high availability.

7. What are the Advantages of Using MongoDB Over Other Databases?

- **Flexibility:** MongoDB's document-oriented model allows for dynamic schemas.
- **Scalability:** Built-in sharding enables horizontal scaling.
- **High Availability:** Replica sets provide redundancy and automatic failover.
- **Performance:** Efficient storage and retrieval with BSON format.
- **Ease of Use:** JSON-like documents make it easier for developers to interact with data.

8. How to Create a New Database and Collection in MongoDB?

To create a new database and collection in MongoDB, you can use the mongo shell:

```
use mydatabase
db.createCollection("mycollection")
```

This command switches to mydatabase (creating it if it doesn't exist) and creates a new collection named mycollection.

9. What is Sharding, and How Does It Work in MongoDB?

Sharding is a method for distributing data across multiple servers in MongoDB. It allows for horizontal scaling by splitting large datasets into smaller, more manageable pieces called shards. Each shard is a separate database that holds a portion of the data. MongoDB automatically balances data and load across shards, ensuring efficient data distribution and high performance.

CRUD operations in MongoDB are used to create, read, update, and delete documents.

- **Create:** `db.collection.insertOne({ name: "Alice", age: 25 })`
- **Read:** `db.collection.find({ name: "Alice" })`
- **Update:** `db.collection.updateOne({ name: "Alice" }, { $set: { age: 26 } })`
- **Delete:** `db.collection.deleteOne({ name: "Alice" })`

11. How to Perform Basic Querying in MongoDB?

Basic querying in MongoDB involves using the `find` method to retrieve documents that match certain criteria. For example:

```
db.collection.find({ age: { $gte: 20 } })
```

This query retrieves all documents from the collection where the `age` field is greater than or equal to 20.

12. What is an Index in MongoDB, and How to Create One?

An index in MongoDB is a data structure that improves the speed of data retrieval operations on a collection. You can create an index using the `createIndex` method. For example, to create an index on the `name` field:

```
db.collection.createIndex({ name: 1 })
```

13. How Does MongoDB Handle Data Consistency?

MongoDB provides several mechanisms to ensure data consistency:

- **Journaling:** MongoDB uses write-ahead logging to maintain data integrity.
- **Write Concerns:** Specify the level of acknowledgment requested from MongoDB for write operations (e.g., acknowledgment from primary only, or acknowledgment from primary and secondaries).
- **Replica Sets:** Replication ensures data is consistent across multiple nodes, and read concerns can be configured to ensure data consistency for read

14. How to Perform Data Import and Export in MongoDB?

To perform data import and export in MongoDB, you can use the `mongoimport` and `mongoexport` tools. These tools allow you to import data from JSON, CSV, or TSV files into MongoDB and export data from MongoDB collections to JSON or CSV files.

Import Data:

```
mongoimport --db mydatabase --collection mycollection --file data.json
```

This command imports data from `data.json` into the `mycollection` collection in the `mydatabase` database.

Export Data:

```
mongoexport --db mydatabase --collection mycollection --out data.json
```

This command exports data from the `mycollection` collection in the `mydatabase` database to `data.json`.

15. What are MongoDB Aggregation Pipelines and How are They Used?

The aggregation pipeline is a framework for data aggregation, modeled on the concept of data processing pipelines. Documents enter a multi-stage pipeline that transforms the documents into aggregated results. Each stage performs an operation on the input documents and passes the results to the next stage.

```
db.orders.aggregate([
  { $match: { status: "A" } }, ..... // Stage 1: Filter documents by
  status
  { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }, // Stage
  2: Group by customer ID and sum the amount
  { $sort: { total: -1 } } ..... // Stage 3: Sort by total in
  descending order
])
```

In this example:

- Stage 2 (\$group) groups documents by customer ID and calculates the total amount for each group.
- Stage 3 (\$sort) sorts the results by total amount in descending order.

Aggregation pipelines are powerful and flexible, enabling complex data processing tasks to be executed within MongoDB.

MongoDB Intermediate Interview Questions

MongoDB Intermediate Interview Questions dive into more complex features of this NoSQL database. These questions might cover how to design a good schema, use different indexing strategies, and work with the aggregation framework.

16. Describe the Aggregation Framework in MongoDB

The Aggregation Framework in MongoDB is a powerful tool for performing data processing and transformation on documents within a collection. It works by passing documents through a multi-stage pipeline, where each stage performs a specific operation on the data, such as filtering, grouping, sorting, reshaping, and computing aggregations. This framework is particularly useful for creating complex data transformations and analytics directly within the database.

17. How to Perform Aggregation Operations Using MongoDB?

Aggregation operations in MongoDB are performed using the `aggregate` method. This method takes an array of pipeline stages, each stage representing a step in the data processing pipeline. For example, to calculate the total sales for each product, you might use the following aggregation pipeline:

```
db.sales.aggregate([
  { $match: { status: "completed" } }, // Filter completed sales
  { $group: { _id: "$product", totalSales: { $sum: "$amount" } } }, //
  Group by product and sum the sales amount
  { $sort: { totalSales: -1 } } // Sort by total sales in descending
```

18. Explain the Concept of Write Concern and Its Importance in MongoDB

Write Concern in MongoDB refers to the level of acknowledgment requested from MongoDB for write operations. It determines how many nodes must confirm the write operation before it is considered successful. Write concern levels range from “acknowledged” (default) to “unacknowledged,” “journaled,” and various “replica acknowledged” levels.

The importance of write concern lies in balancing between data durability and performance. Higher write concern ensures data is safely written to disk and replicated, but it may impact performance due to the added latency.

19. What are TTL Indexes, and How are They Used in MongoDB?

TTL (Time To Live) Indexes in MongoDB are special indexes that automatically remove documents from a collection after a certain period. They are commonly used for data that needs to expire after a specific time, such as session information, logs, or temporary data. To create a TTL index, you can specify the expiration time in seconds:

```
db.sessions.createIndex({ "createdAt": 1 }, { expireAfterSeconds: 3600 })
```

This index will remove documents from the sessions collection 1 hour (3600 seconds) after the createdAt field's value.

20. How to Handle Schema Design and Data Modeling in MongoDB?

Schema design and data modeling in MongoDB involve defining how data is organized and stored in a document-oriented database. Unlike SQL databases, MongoDB offers flexible schema design, which can be both an advantage and a challenge. Key considerations for schema design include:

- **Embedding vs. Referencing:** Deciding whether to embed related data within a single document or use references between documents.
- **Document Structure:** Designing documents that align with application query

- **Data Duplication:** Accepting some level of data duplication to optimize for read performance.
- **Sharding:** Designing the schema to support sharding if horizontal scaling is required.

21. What is GridFS, and When is it Used in MongoDB?

GridFS is a specification for storing and retrieving large files in MongoDB. It is used when files exceed the BSON-document size limit of 16 MB or when you need to perform efficient retrieval of specific file sections. GridFS splits a large file into smaller chunks and stores each chunk as a separate document within two collections: `fs.files` and `fs.chunks`. This allows for efficient storage and retrieval of large files, such as images, videos, or large datasets.

22. Explain the Differences Between WiredTiger and MMAPv1 Storage Engines

WiredTiger and MMAPv1 are two different storage engines in MongoDB, each with distinct characteristics:

WiredTiger:

- **Concurrency:** Provides document-level concurrency, allowing multiple operations to be performed simultaneously.
- **Compression:** Supports data compression, reducing storage space requirements.
- **Performance:** Generally offers better performance and efficiency for most workloads.
- **Journaling:** Uses a write-ahead logging mechanism for better data integrity.

MMAPv1:

- **Concurrency:** Provides collection-level concurrency, which can limit performance under heavy write operations.
- **No Compression:** Does not support data compression.
- **Legacy Engine:** It is the older storage engine and has been deprecated in

- **Simplicity:** Simple implementation but lacks advanced features compared to WiredTiger.

23. How to Handle Transactions in MongoDB?

MongoDB supports multi-document ACID transactions, allowing you to perform a series of read and write operations across multiple documents and collections in a transaction. This ensures data consistency and integrity. To use transactions, you typically start a session, begin a transaction, perform the operations, and then commit or abort the transaction.

Example in JavaScript:

```
const session = client.startSession();

session.startTransaction();

try {
  db.collection1.insertOne({ name: "Alice" }, { session });
  db.collection2.insertOne({ name: "Bob" }, { session });
  session.commitTransaction();
} catch (error) {
  session.abortTransaction();
} finally {
  session.endSession();
}
```

24. Describe the MongoDB Compass Tool and Its Functionalities

MongoDB Compass is a [graphical user interface](#) (GUI) tool for MongoDB that provides an easy way to visualize, explore, and manipulate your data. It offers features such as:

- **Schema Visualization:** View and analyze your data schema, including field types and distributions.
- **Query Building:** Build and execute queries using a visual interface.
- **Aggregation Pipeline:** Construct and run aggregation pipelines.

- **Index Management:** Create and manage indexes to optimize query performance.
- **Performance Monitoring:** Monitor database performance, including slow queries and resource utilization.
- **Data Validation:** Define and enforce schema validation rules to ensure data integrity.
- **Data Import/Export:** Easily import and export data between MongoDB and JSON/CSV files.

25. What is MongoDB Atlas, and How Does it Differ From Self-Hosted MongoDB?

MongoDB Atlas is a fully managed cloud database service provided by MongoDB. It offers automated deployment, scaling, and management of MongoDB clusters across various cloud providers ([AWS](#), [Azure](#), [Google Cloud](#)). Key differences from self-hosted MongoDB include:

- **Managed Service:** Atlas handles infrastructure management, backups, monitoring, and upgrades.
- **Scalability:** Easily scale clusters up or down based on demand.
- **Security:** Built-in security features such as encryption, access controls, and compliance certifications.
- **Global Distribution:** Deploy clusters across multiple regions for low-latency access and high availability.
- **Integrations:** Seamless integration with other cloud services and MongoDB tools.

26. How to Implement Access Control and User Authentication in MongoDB?

Access control and user authentication in MongoDB are implemented through a role-based access control (RBAC) system. You create users and assign roles that define their permissions. To set up access control:

- **Enable Authentication:** Configure MongoDB to require authentication by

- **Create Users:** Use the `db.createUser` method to create users with specific roles.

```
db.createUser({
  user: "admin",
  pwd: "password",
  roles: [{ role: "userAdminAnyDatabase", db: "admin" }]
});
```

- **Assign Roles:** Assign roles to users that define their permissions, such as read, write, or admin roles for specific databases or collections.

27. What are Capped Collections, and When are They Useful?

Capped collections in MongoDB are fixed-size collections that automatically overwrite the oldest documents when the specified size limit is reached. They maintain insertion order and are useful for scenarios where you need to store a fixed amount of recent data, such as logging, caching, or monitoring data.

Example of creating a capped collection:

```
db.createCollection("logs", { capped: true, size: 100000 });
```

28. Explain the Concept of Geospatial Indexes in MongoDB

Geospatial indexes in MongoDB are special indexes that support querying of geospatial data, such as locations and coordinates. They enable efficient queries for proximity, intersections, and other spatial relationships. MongoDB supports two types of geospatial indexes: 2d for flat geometries and 2dsphere for spherical geometries.

Example of creating a 2dsphere index:

```
db.places.createIndex({ location: "2dsphere" });
```

29. How to Handle Backups and Disaster Recovery in MongoDB?

Handling backups and disaster recovery in MongoDB involves regularly creating backups of your data and having a plan for restoring data in case of failure. Methods include:

- **Mongodump/Mongorestore:** Use the mongodump and mongorestore utilities to create and restore binary backups.
- **File System Snapshots:** Use file system snapshots to take consistent backups of the data files.
- **Cloud Backups:** If using MongoDB Atlas, leverage automated backups provided by the service.
- **Replica Sets:** Use replica sets to ensure data redundancy and high availability. Regularly test the failover and recovery process.

30. Describe the Process of Upgrading MongoDB to a Newer Version

Upgrading MongoDB to a newer version involves several steps to ensure a smooth transition:

- **Check Compatibility:** Review the release notes and compatibility changes for the new version.
- **Backup Data:** Create a backup of your data to prevent data loss.
- **Upgrade Drivers:** Ensure that your application drivers are compatible with the new MongoDB version.
- **Upgrade MongoDB:** Follow the official MongoDB upgrade instructions, which typically involve stopping the server, installing the new version, and restarting the server.
- **Test Application:** Thoroughly test your application with the new MongoDB version to identify any issues.
- **Monitor:** Monitor the database performance and logs to ensure a successful upgrade.

31. What are Change Streams in MongoDB, and How are They Used?

Change Streams in MongoDB allow applications to listen for real-time changes to data in collections, databases, or entire clusters. They provide a powerful

replace, and delete operations. To use Change Streams, you typically open a change stream cursor and process the change events as they occur.

Example:

```
const changeStream = db.collection('orders').watch();
changeStream.on('change', (change) => {
  console.log(change);
});
```

This example listens for changes in the orders collection and logs the change events.

32. Explain the Use of Hashed Sharding Keys in MongoDB

Hashed Sharding Keys in MongoDB distribute data across shards using a hashed value of the shard key field. This approach ensures an even distribution of data and avoids issues related to data locality or uneven data distribution that can occur with range-based sharding. Hashed sharding is useful for fields with monotonically increasing values, such as timestamps or identifiers.

Example:

```
db.collection.createIndex({ _id: "hashed" });
sh.shardCollection("mydb.mycollection", { _id: "hashed" });
```

33. How to Optimize MongoDB Queries for Performance?

Optimizing MongoDB queries involves several strategies:

- **Indexes:** Create appropriate indexes to support query patterns.
- **Query Projections:** Use projections to return only necessary fields.
- **Index Hinting:** Use index hints to force the query optimizer to use a specific index.
- **Query Analysis:** Use the explain() method to analyze query execution plans and identify bottlenecks.
- **Aggregation Pipeline:** Optimize the aggregation pipeline stages to minimize

34. Describe the Map-Reduce Functionality in MongoDB

Map-Reduce in MongoDB is a data processing paradigm used to perform complex data aggregation operations. It consists of two phases: the map phase processes each input document and emits key-value pairs, and the reduce phase processes all emitted values for each key and outputs the final result.

Example:

```
db.collection.mapReduce(  
  function() { emit(this.category, this.price); },  
  function(key, values) { return Array.sum(values); },  
  { out: "category_totals" }  
);
```

This example calculates the total price for each category in a collection.

35. What is the Role of Journaling in MongoDB, and How Does It Impact Performance?

Journaling in MongoDB ensures data durability and crash recovery by recording changes to the data in a journal file before applying them to the database files. This mechanism allows MongoDB to recover from unexpected shutdowns or crashes by replaying the journal. While journaling provides data safety, it can impact performance due to the additional I/O operations required to write to the journal file.

36. How to Implement Full-Text Search in MongoDB?

Full-Text Search in MongoDB is implemented using text indexes. These indexes allow you to perform text search queries on string content within documents.

Example:

```
db.collection.createIndex({ content: "text" });  
db.collection.find({ $text: { $search: "mongodb" } }):
```

In this example, a text index is created on the content field, and a text search query is performed to find documents containing the word “mongodb.”

37. What are the Considerations for Deploying MongoDB in a Production Environment?

Considerations for deploying MongoDB in a production environment include:

- Replication: Set up replica sets for high availability and data redundancy.
- Sharding: Implement sharding for horizontal scaling and to distribute the load.
- Backup and Recovery: Establish a robust backup and recovery strategy.
- Security: Implement authentication, authorization, and encryption.
- Monitoring: Use monitoring tools to track performance and detect issues.
- Capacity Planning: Plan for adequate storage, memory, and CPU resources.
- Maintenance: Regularly update MongoDB to the latest stable version and perform routine maintenance tasks.

38. Explain the Concept of Horizontal Scalability and Its Implementation in MongoDB

Horizontal Scalability in MongoDB refers to the ability to add more servers to distribute the load and data. This is achieved through sharding, where data is partitioned across multiple shards. Each shard is a replica set that holds a subset of the data. Sharding allows MongoDB to handle large datasets and high-throughput operations by distributing the workload.

39. How to Monitor and Troubleshoot Performance Issues in MongoDB?

Monitoring and troubleshooting performance issues in MongoDB involve:

- Monitoring Tools: Use tools like MongoDB Cloud Manager, MongoDB Ops Manager, or third-party monitoring solutions.
- Logs: Analyze MongoDB logs for errors and performance metrics.
- Profiling: Enable database profiling to capture detailed information about

- **Explain Plans:** Use the `explain()` method to understand query execution and identify bottlenecks.
- **Index Analysis:** Review and optimize indexes based on query patterns and usage.
- **Resource Utilization:** Monitor CPU, memory, and disk I/O usage to identify resource constraints.

40. Describe the Process of Migrating Data from a Relational Database to MongoDB

Migrating data from a relational database to MongoDB involves several steps:

- **Schema Design:** Redesign the relational schema to fit MongoDB's document-oriented model. Decide on embedding vs. referencing, and plan for indexes and collections.
- **Data Export:** Export data from the relational database in a format suitable for MongoDB (e.g., CSV, JSON).
- **Data Transformation:** Transform the data to match the MongoDB schema. This can involve converting data types, restructuring documents, and handling relationships.
- **Data Import:** Import the transformed data into MongoDB using tools like `mongoimport` or custom scripts.
- **Validation:** Validate the imported data to ensure consistency and completeness.
- **Application Changes:** Update the application code to interact with MongoDB instead of the relational database.
- **Testing:** Thoroughly test the application and the database to ensure everything works as expected.
- **Go Live:** Deploy the MongoDB database in production and monitor the transition.

MongoDB Query Based Interview Questions

MongoDB Query-Based Interview Questions focus on your ability to write and optimize queries to interact with the database effectively. These questions

```
"[
  {
    "_id": 1,
    "name": "John Doe",
    "age": 28,
    "position": "Software Engineer",
    "salary": 80000,
    "department": "Engineering",
    "hire_date": ISODate("2021-01-15")
  },
  {
    "_id": 2,
    "name": "Jane Smith",
    "age": 34,
    "position": "Project Manager",
    "salary": 95000,
    "department": "Engineering",
    "hire_date": ISODate("2019-06-23")
  },
  {
    "_id": 3,
    "name": "Emily Johnson",
    "age": 41,
    "position": "CTO",
    "salary": 150000,
    "department": "Management",
    "hire_date": ISODate("2015-03-12")
  },
  {
    "_id": 4,
    "name": "Michael Brown",
    "age": 29,
    "position": "Software Engineer",
    "salary": 85000,
    "department": "Engineering",
    "hire_date": ISODate("2020-07-30")
  },
  {

```



```
.....  ""age"": 26,
.....  ""position"": ""UI/UX Designer"",
.....  ""salary"": 70000,
.....  ""department"": ""Design"",
.....  ""hire_date"": ISODate("""2022-10-12""")
.....  }
..... ]"
```

1. Find all Employees Who Work in the “Engineering” Department.

Query:

```
db.employees.find({ department: "Engineering" })
```

Output:

```
[
  {
    "_id": 1,
    "name": "John Doe",
    "age": 28,
    "position": "Software Engineer",
    "salary": 80000,
    "department": "Engineering",
    "hire_date": ISODate("2021-01-15")
  },
  {
    "_id": 2,
    "name": "Jane Smith",
    "age": 34,
    "position": "Project Manager",
    "salary": 95000,
    "department": "Engineering",
    "hire_date": ISODate("2019-06-23")
  },
  {
    "_id": 4,
    "name": "Michael Brown".
```

```
..... "salary": 85000,  
..... "department": "Engineering",  
..... "hire_date": ISODate("2020-07-30")  
..... }  
..... ]
```

Explanation:

This query finds all employees whose department field is “Engineering”.

2. Find the Employee with the Highest Salary.**Query:**

```
db.employees.find().sort({ salary: -1 }).limit(1)
```

Output:

```
[  
  {  
    "_id": 3,  
    "name": "Emily Johnson",  
    "age": 41,  
    "position": "CTO",  
    "salary": 150000,  
    "department": "Management",  
    "hire_date": ISODate("2015-03-12")  
  }  
]
```

Explanation:

This query sorts all employees by salary in descending order and retrieves the top document, which is the employee with the highest salary.

3. Update the Salary of “John Doe” to 90000.**Query:**

Output:

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" : 1 }
```

Explanation:

This query updates the salary of the employee named “John Doe” to 90000.

4. Count the Number of Employees in Each Department.**Query:**

```
db.employees.aggregate([
  { $group: { _id: "$department", count: { $sum: 1 } } }
])
```

Output:

```
[
  { "_id": "Engineering", "count": 3 },
  { "_id": "Management", "count": 1 },
  { "_id": "Design", "count": 1 }
]
```

Explanation:

This query groups the employees by the department field and counts the number of employees in each department.

5. Add a New Field Bonus to All Employees in the “Engineering” Department with a Value of 5000.**Query:**

```
db.employees.updateMany({ department: "Engineering" }, { $set: { bonus: 5000 } })
```

Output:

Explanation:

This query adds a new field bonus with a value of 5000 to all employees in the “Engineering” department.

6. Retrieve All Documents in the Employees Collection and Sort Them by the Length of Their Name in Descending Order.**Query:**

```
db.employees.aggregate([
  { $addFields: { nameLength: { $strLenCP: "$name" } } },
  { $sort: { nameLength: -1 } },
  { $project: { nameLength: 0 } }
])
```

Output:

```
[
  {
    "_id": 2,
    "name": "Jane Smith",
    "age": 34,
    "position": "Project Manager",
    "salary": 95000,
    "department": "Engineering",
    "hire_date": ISODate("2019-06-23")
  },
  {
    "_id": 3,
    "name": "Emily Johnson",
    "age": 41,
    "position": "CTO",
    "salary": 150000,
    "department": "Management",
    "hire_date": ISODate("2015-03-12")
  },
  {

```

```
..... "age": 28,
..... "position": "Software Engineer",
..... "salary": 80000,
..... "department": "Engineering",
..... "hire_date": ISODate("2021-01-15")
..... },
..... {
.....   "_id": 4,
.....   "name": "Michael Brown",
.....   "age": 29,
.....   "position": "Software Engineer",
.....   "salary": 85000,
.....   "department": "Engineering",
.....   "hire_date": ISODate("2020-07-30")
..... },
..... {
.....   "_id": 5,
.....   "name": "Sarah Davis",
.....   "age": 26,
.....   "position": "UI/UX Designer",
.....   "salary": 70000,
.....   "department": "Design",
.....   "hire_date": ISODate("2022-10-12")
..... }
..... ]
```

Explanation:

This query calculates the length of each employee's name, sorts the documents by this length in descending order, and removes the temporary nameLength field from the output.

7. Find the Average Salary of Employees in the “Engineering” Department.

Query:

```
db.employees.aggregate([
  { $match: { department: "Engineering" } },
```

Output:

```
[
  { "_id": null, "averageSalary": 86666.66666666667 }
]
```

Explanation:

This query filters employees to those in the “Engineering” department and calculates the average salary of these employees.

8. Find the Department with the Highest Average Salary.**Query:**

```
db.employees.aggregate([
  { $group: { _id: "$department", averageSalary: { $avg: "$salary" } } },
  { $sort: { averageSalary: -1 } },
  { $limit: 1 }
])
```

Output:

```
[
  { "_id": "Management", "averageSalary": 150000 }
]
```

Explanation:

This query groups employees by department, calculates the average salary for each department, sorts these averages in descending order, and retrieves the department with the highest average salary.

9. Find the Total Number of Employees Hired in Each Year.**Query:**

```
}  
])
```

Output:

```
[  
  { "_id": 2015, "totalHired": 1 },  
  { "_id": 2019, "totalHired": 1 },  
  { "_id": 2020, "totalHired": 1 },  
  { "_id": 2021, "totalHired": 1 },  
  { "_id": 2022, "totalHired": 1 }  
]
```

Explanation:

This query groups employees by the year they were hired, which is extracted from the hire_date field, and counts the total number of employees hired each year.

10. Find the Highest and Lowest Salary in the “Engineering” Department.**Query:**

```
db.employees.aggregate([  
  { $match: { department: "Engineering" } },  
  {  
    $group: {  
      _id: null,  
      highestSalary: { $max: "$salary" },  
      lowestSalary: { $min: "$salary" }  
    }  
  }  
])
```

Output:

```
[  
  { "_id": null, "highestSalary": 95000, "lowestSalary": 80000 }  
]
```


This query filters employees to those in the “Engineering” department, then calculates the highest and lowest salary within this group.

Conclusion

Preparing for a MongoDB interview can be much easier with the right set of questions. This blog post has provided a comprehensive collection of MongoDB interview questions, ranging from basic to advanced topics. By practicing these questions, you can build your confidence and improve your understanding of key MongoDB concepts.

v vivek... + Follow



1

Next Article

Top 50 Database Interview Questions and
Answers for 2024

Similar Reads

Top 50 Database Interview Questions and Answers for 2024

Getting ready for a database interview can be challenging, but knowing the types of questions that you can encounter in the interview can help you feel more...

15+ min read

Top 15 Uber SQL Interview Questions with Answers and Explanations

Preparing for a SQL interview at Uber involves not only understanding SQL concepts but also being able to work with practical examples. Here, we'll provid...

12 min read

To 50 Oracle Interview Questions and Answers for 2024

When you are preparing for any Interview, the big question is what kind of questions they'll ask during the interview. Now to make this In this blog post

Best 50 PostgreSQL Interview Questions and Answers for 2024

Discover the ultimate resource for PostgreSQL interview preparation with our curated list of high-impact questions. Whether you're gearing up for a...

15+ min read

How to Integrate MongoDB Atlas and Segment using MongoDB Stitch

Data integration has become a crucial component of modern business strategies, allowing companies to enhance the power of data for informed decision-making....

5 min read

Connect MongoDB Atlas Cluster to MongoDB Compass

MongoDB Compass is a free GUI for MongoDB. You might want to connect MongoDB Atlas Cluster to MongoDB Compass to take benefit of the GUI model...

4 min read

MongoDB Shell vs MongoDB Node.JS Driver

MongoDB Shell and the MongoDB Node.js Driver are both tools for interacting with MongoDB, but they serve different purposes and are used in different...

4 min read

MongoDB Compass vs MongoDB Atlas

MongoDB is a popular NoSQL database known for flexibility and scalability. MongoDB Compass offers a graphical interface for managing databases visually...

5 min read

Top 7 MongoDB Tools for 2024

MongoDB is one of the most popular databases in the world, as it is used by many big organizations. It provides a convenient way to store data. But to use...

9 min read

PL/SQL Interview Questions and Answers

PL/SQL (Procedural Language for SQL) is one of the most widely used database programming languages in the world, known for its robust performance....

Top 50 SQLite Interview Questions for 2024

In this interview preparation blog post, we'll explore some of the most common SQLite interview questions that can help you showcase your knowledge and...

15+ min read

Top 10 MongoDB Projects Ideas for Beginners

Due to the increase in demand for the latest technologies, different databases are used in multiple industries. MongoDB is one of the well-known non-relational...

9 min read

MongoDB vs Orient DB: Top Differences

The appropriate selection of database technology can determine successful application development. About NoSQL databases, MongoDB and OrientDB are...

8 min read

Connecting MongoDB to Jupyter Notebook

MongoDB is one of the most famous NoSql databases. It is an open-source database. The simple meaning of the NoSql database is a non-relational...

3 min read

Create a Newsletter Sourcing Data using MongoDB

There are many news delivery websites available like ndtv.com. In this article, let us see the very useful and interesting feature of how to get the data from...

7 min read

How to get Distinct Documents from MongoDB using Node.js ?

MongoDB is a cross-platform, document-oriented database that works on the concept of collections and documents. It stores data in the form of key-value pai...

2 min read

MongoDB - Current Date Operator (\$currentDate)

MongoDB provides different types of field update operators to update the values of the fields of the documents and \$currentDate operator is one of them. This...

3D Plotting sample Data from MongoDB Atlas Using Python

MongoDB, the most popular NoSQL database, is an open-source document-oriented database. The term 'NoSQL' means 'non-relational'. It means that...

3 min read

Native MongoDB driver for Node.js

The native MongoDB driver for Node.JS is a dependency that allows our JavaScript application to interact with the NoSQL database, either locally or on...

5 min read

Python MongoDB - Update_many Query

MongoDB is a NoSQL database management system. Unlike MySQL the data in MongoDB is not stored as relations or tables. Data in mongoDB is stored as...

3 min read

MongoDB | ObjectID() Function

ObjectID() Function: MongoDB uses ObjectID to create unique identifiers for all the documents in the database. It is different than the traditional...

2 min read

RESTfull routes on Node.js and MongoDB

REST stands for representational state transfer which basically provides a way of mapping HTTP verbs such as (GET, POST, PUT, DELETE) with CRUD actions suc...

6 min read

Python MongoDB - \$group (aggregation)

MongoDB is an open-source document-oriented database. MongoDB stores data in the form of key-value pairs and is a NoSQL database program. The term...

3 min read

Python MongoDB - distinct()

MongoDB is a cross-platform, document-oriented database that works on the concept of collections and documents. It stores data in the form of key-value pai...

Difference between Impala and MongoDB

1. Impala : Impala is a query engine that runs on Hadoop. It is an open source software and massively parallel processing SQL query engine. It supports in-...

2 min read

Difference between MongoDB and ActivePivot

1. MongoDB : MongoDB is an open-source document-oriented database used for high volume data storage. It falls under classification of NoSQL database. NoSQ...

2 min read

Difference between dBASE and MongoDB

1. dBASE : dBASE was one of the most successful database management systems for microcomputers. It was the first commercially successful database...

2 min read

MongoDB \$In Operator

MongoDB provides different types of arithmetic expression operators that are used in the aggregation pipeline stages and \$In operator is one of them. This...

1 min read

MongoDB \$sqrt Operator

MongoDB provides different types of arithmetic expression operators that are used in the aggregation pipeline stages \$sqrt operator is one of them. This...

2 min read

MongoDB \$mod Operator

MongoDB provides different types of arithmetic expression operators that are used in the aggregation pipeline stages and \$mod operator is one of them. This...

1 min read

Article Tags :

[MongoDB](#)

[Databases](#)



Corporate & Communications Address:-
A-143, 9th Floor, Sovereign Corporate
Tower, Sector- 136, Noida, Uttar Pradesh
(201305) | Registered Address:- K 061,
Tower K, Gulshan Vivante Apartment,
Sector 137, Noida, Gautam Buddh
Nagar, Uttar Pradesh, 201305



Company

About Us
Legal
Careers
In Media
Contact Us
Advertise with us
GFG Corporate Solution
Placement Training Program

Languages

Python
Java
C++
PHP

Explore

Job-A-Thon Hiring Challenge
Hack-A-Thon
GfG Weekly Contest
Offline Classes (Delhi/NCR)
DSA in JAVA/C++
Master System Design
Master CP
GeeksforGeeks Videos
Geeks Community

DSA

Data Structures
Algorithms
DSA for Beginners
Basic DSA Problems

Android Tutorial

Data Science & ML

Data Science With Python

Data Science For Beginner

Machine Learning

ML Maths

Data Visualisation

Pandas

NumPy

NLP

Deep Learning

Python Tutorial

Python Programming Examples

Django Tutorial

Python Projects

Python Tkinter

Web Scraping

OpenCV Tutorial

Python Interview Question

DevOps

Git

AWS

Docker

Kubernetes

Azure

GCP

DevOps Roadmap

School Subjects

Mathematics

Physics

Chemistry

Biology

Social Science

English Grammar

Web Technologies

HTML

CSS

JavaScript

TypeScript

ReactJS

NextJS

NodeJs

Bootstrap

Tailwind CSS

Computer Science

GATE CS Notes

Operating Systems

Computer Network

Database Management System

Software Engineering

Digital Logic Design

Engineering Maths

System Design

High Level Design

Low Level Design

UML Diagrams

Interview Guide

Design Patterns

OOAD

System Design Bootcamp

Interview Questions

Commerce

Accountancy

Business Studies

Economics

Management

HR Management

Finance

Income Tax

Databases

SQL
MYSQL
PostgreSQL
PL/SQL
MongoDB

Preparation Corner

Company-Wise Recruitment Process
Resume Templates
Aptitude Preparation
Puzzles
Company-Wise Preparation
Companies
Colleges

Competitive Exams

JEE Advanced
UGC NET
UPSC
SSC CGL
SBI PO
SBI Clerk
IBPS PO
IBPS Clerk

More Tutorials

Software Development
Software Testing
Product Management
Project Management
Linux
Excel
All Cheat Sheets
Recent Articles

Free Online Tools

Typing Test
Image Editor
Code Formatters
Code Converters
Currency Converter
Random Number Generator
Random Password Generator

Write & Earn

Write an Article
Improve an Article
Pick Topics to Write
Share your Experiences
Internships

@GeeksforGeeks, Sanchhaya Education Private Limited, All rights reserved